

1)

$$\begin{aligned} \text{a) Cost of accepting} &= P[y = +1|x] * 0 + P[y = -1|x] * c_a \\ &= (1 - g(x))c_a \\ \text{Cost of rejecting} &= P[y = +1|x] * c_r + P[y = -1|x] * 0 \\ &= g(x)c_r \end{aligned}$$

b) We only want to accept when the cost of accepting is  $\leq$  the cost of rejecting.  
Plugging in from what we derived in part a :

$$\begin{aligned} &\rightarrow (1 - g(x))c_a \leq g(x)c_r \\ &\rightarrow c_a - c_a * g(x) \leq g(x)c_r \\ &\rightarrow c_a \leq g(x)(c_a + c_r) \\ &\rightarrow \frac{c_a}{c_a + c_r} \leq g(x) \end{aligned}$$

The Threshold =  $\frac{c_a}{c_a + c_r}$  hence we reject when  $\frac{c_a}{c_a + c_r} > g(x)$

$$\text{c) Supermarket threshold} = \frac{1}{10+1} = \frac{1}{11}$$

This threshold accommodates the fact that we would want to avoid a false negative more than a false positive.

$$\text{CIA threshold} = \frac{1}{1000+1} = \frac{1}{1001}$$

In this case, we want to avoid false positives since letting in a dangerous acceptance can be detrimental. Hence, the high threshold makes sense.

2)

## 2a

```
In [1]: 1 from sklearn.linear_model import LogisticRegression
2 from sklearn.datasets import load_wine
3 from sklearn.model_selection import cross_val_score
4 import numpy as np
5 import matplotlib.pyplot as plt
6 import math
```

```
In [2]: 1 def makePoly(deg):
2     coefs = np.random.normal(size=deg+1)
3     coefs = np.true_divide(coefs, np.sqrt(np.sum(coefs**2)))
4     return coefs
```


```
In [3]: 1 def randSample(numPoints, p, noiseStd):
2     xi = np.random.uniform(-1, 1, size=numPoints)
3     noise = np.random.normal(scale=noiseStd, size=numPoints)
4     p_xi = p(xi + noise)
5     return xi, p_xi
```

```
In [4]: 1 def oneFit(numPoints=10, degTarget=4, degFit=2, noiseStd=0.1):
2     coefs = makePoly(degTarget)
3     p = np.poly1d(coefs)
4     xi, pred_xi = randSample(numPoints, p, noiseStd)
5     #target function
6     plot(p)
7     #points
8     plt.scatter(xi, pred_xi)
9     #predicted graph
10    p_prime = np.poly1d(np.polyfit(xi, pred_xi, degFit))
11    plot(p_prime)
12    #calc e_out
13    x = np.arange(-2, 2, 0.05)
14    y = p(x)
15    y_pred = p_prime(x)
16    error_out = 0
17    for i in range(len(y)):
18        error_out += (y[i] - y_pred[i])**2
19    error_out = error_out/len(y)
20    #calc e_in
21    error_in = 0
22    for j in range(len(xi)):
23        error_in += (p_prime(xi[j]) - pred_xi[j])**2
24    error_in = error_in/len(xi)
25    return error_in, error_out
26
```

```
In [5]: 1 def plot(p):
2     x = np.arange(-2, 2, 0.05)
3     y = p(x)
4     plt.plot(x, y)
```

```
In [6]: 1 oneFit()
```

Out[6]: (0.004506580022704715, 2.610329133969736)




## 2b





```
In [7]: 1 #no Plot
2 def oneFit_noPrint(numPoints=10, degTarget=4, degFit=2, noiseStd=0.1):
3     coefs = makePoly(degTarget)
4     p = np.polyd(coefs)
5     xi, pred_xi = randSample(numPoints, p, noiseStd)
6     p_prime = np.polyd(np.polyfit(xi, pred_xi, degFit))
7     #calc e_out
8     x = np.arange(-2, 2, 0.05)
9     y = p(x)
10    y_pred = p_prime(x)
11    error_out = 0
12    for i in range(len(y)):
13        error_out += (y[i] - y_pred[i])**2
14    error_out = error_out/len(y)
15    #calc e_in
16    error_in = 0
17    for j in range(len(xi)):
18        error_in += (p_prime(xi[j]) - pred_xi[j])**2
19    error_in = error_in/len(xi)
20    return error_in, error_out






In [8]: 1 def testRangeDegrees(degrees=6, sampleSize=20, degTarget=4, noiseStd=0.1):
2     answer = np.empty((degrees, 2))
3     for counter in range(degrees):
4         counter_2 = counter + 1
5         b_avg = np.empty((1000, 2))
6         for i in range(1000):
7             exp_in, exp_out = oneFit_noPrint(sampleSize, degTarget, counter_2, noiseStd)
8             b_avg[i, 0] = exp_in
9             b_avg[i, 1] = exp_out
10        answer[counter] = np.mean(b_avg, axis=0)
11    return answer
12
13 a = testRangeDegrees()
14 a
15
```


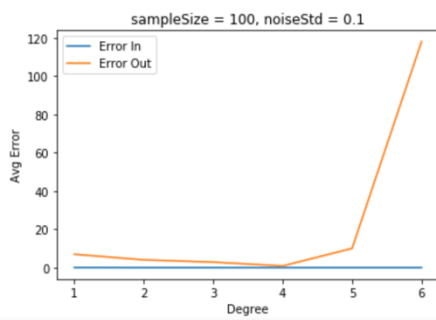



```
Out[8]: array([[4.44087945e-02, 6.96669436e+00],
               [1.52368202e-02, 4.27435777e+00],
               [9.97849182e-03, 3.70950625e+00],
               [8.73317273e-03, 6.32141119e+00],
               [7.55428524e-03, 1.03322615e+02],
               [7.53490191e-03, 5.01167410e+03]])
```


```
In [9]: 1 samples = [10, 20, 50, 100]
2 noise = [0, 0.1, 0.2, 0.5]
3 x_axis = [1, 2, 3, 4, 5, 6]
4
5 for size in samples:
6     for n in noise:
7         avg_error = testRangeDegrees(degrees = 6, sampleSize= size, degTarget= 4, noiseStd=n)
8         plt.figure()
9         plt.plot(x_axis, avg_error[:, [0]], label = 'Error In')
10        plt.plot(x_axis, avg_error[:, [1]], label = 'Error Out')
11        plt.ylabel("Avg Error")
12        plt.xlabel("Degree")
13        plt.title("sampleSize = {}, noiseStd = {}".format(size,n))
14        plt.legend()
15        plt.show()
16
```











When noise is = 0, the out of sample error is minimized regardless of sample


3)

### 3a

```
In [10]: 1 wine = load_wine()
2 data = wine.data
3 target = wine.target
```

```
In [11]: 1 n = list(np.arange(-8,9))
2 x_axis = []
3 acc = []
4 for element in n:
5     c = 2**(0.5*element)
6     x_axis.append(math.log2(c))
7     clf = LogisticRegression(solver = 'liblinear', C = c)
8     clf.fit(data, target)
9     score = cross_val_score(clf, data, target, cv=4)
10    acc.append(np.mean(score))
```

```
In [12]: 1 plt.xlabel('Regularization Constant')
2 plt.ylabel('Accuracy')
3 plt.plot(x_axis, acc)
4 plt.show()
```



### 3b

There looks to be a sweet spot in which the accuracy peaks. This makes sense since the regularization constant, if too big, it would start to block the model from accurately predicting the data. If too small, the noise would still have too much of an effect causing overfitting to occur. The ultimate goal of regularization, when coupled with validation, is to find a constant such that the effect of noise is dimmed while not hurting the model's ability to accurately train on data

4)

## 4a

1602

Leave one out validation means that one sample is not used when validating. The number of samples in the Wine data set is 178, which means there are 178 C 1 ways to pick the sample left out. Since there are three p values and three c values, it makes since that we must train the classifiers 178 C 1 \* 3 \* 3 times = 1602

## 4b

```
In [13]: 1 from sklearn.model_selection import GridSearchCV, LeaveOneOut
        2 from sklearn.model_selection import train_test_split
        3 from sklearn.svm import SVC
        4 from sklearn.metrics import classification_report
```

```
In [14]: 1 numSamples = len(wine.data)
        2 numSamples
```

Out[14]: 178

```
In [15]: 1 tuned_parameters = [{'kernel': ['poly'], 'C': [0.5, 1, 2], 'degree': [1,2,4]}]
        2 cross_val = LeaveOneOut()
        3
        4
```

```
In [18]: 1
        2 clf = GridSearchCV(SVC(gamma = 'scale'), tuned_parameters, iid = False)
        3 clf.fit(data, target)
        4
        5 print("Best parameters is:")
        6 print(clf.best_params_)
        7
```

/Users/alexanderxu/anaconda3/lib/python3.7/site-packages/sklearn/model\_selection/\_split.py:2053: FutureWarning: You should specify a value for 'cv' instead of relying on the default value. The default value will change from 3 to 5 in version 0.22.

warnings.warn(CV\_WARNING, FutureWarning)

Best parameters is:  
{'C': 1, 'degree': 2, 'kernel': 'poly'}

In [ ]:

1



5)

a)  $\nabla E_{\text{aug}}(w(t)) = \nabla E_{\text{in}}(w(t)) + 2\lambda w(t)$

Using this, we can find that:

$$\begin{aligned} w(t+1) &\leftarrow w(t) - \eta \nabla E_{\text{aug}}(w(t)) \\ &= w(t) - \eta (\nabla E_{\text{in}}(w(t)) + 2\lambda w(t)) \\ &= w(t) - \eta \nabla E_{\text{in}}(w(t)) + 2\eta\lambda w(t) \\ &= \mathbf{(1-2\eta\lambda)w(t) - \eta \nabla E_{\text{in}}(w(t))} \end{aligned}$$

b) Given  $y_i \neq w^T x_i$

$$e_i(w) = \max(0, 1 - y_i w^T x_i)$$

Hence:

$$\nabla e_i(w) = \begin{cases} -y_i x_i, & h_i(w) \leq 0 \\ 0, & h_i(w) > 0 \end{cases}$$

c) Using the logic from above:

$$\nabla e_i(w) = \begin{cases} -y_i x_i, & h_i(w) \leq 0 \\ 0, & h_i(w) > 0 \end{cases}$$

But we can substitute  $h_i(w)$  for  $1 - y_i w^T x_i$  thus getting:

$$\nabla e_i(w) = \begin{cases} -y_i x_i, & 1 - y_i w^T x_i \leq 0 \\ 0, & 1 - y_i w^T x_i > 0 \end{cases}$$

Thus, if  $1 - y_i w^T x_i \leq 0$ , meaning we need to update, then:

From part a we can sub:

$$w(t+1) = (1-2\eta\lambda)w(t) - \eta \nabla E_{\text{in}}(w(t))$$

And further make this into:

$$\begin{aligned} w(t+1) &= (1-2\eta\lambda)(1 - y_i w^T x_i) - \eta \nabla E_{\text{in}}(w(t)) \\ w(t+1) &= 1 - y_i w^T x_i - 2\eta\lambda + 2\eta\lambda y_i w^T x_i \\ w(t+1) &= 1 - y_i w^T x_i - 2\eta\lambda(1 - y_i w^T x_i) \\ w(t+1) &= -(y_i w^T x_i - 1) - 2\eta\lambda(1 - y_i w^T x_i) \end{aligned}$$

If  $1 - y_i w^T x_i > 0$  then we do nothing since no update is needed